

Project 3

Time Series Forecasting - Stock Price Prediction

Description for Stock Market Prices Data and Its Features:

Stock market data is financial market data relating to the historical, real time and future values of equities.

- **Open** - The price the stock opened at.
- **High** - The highest price during the day.
- **Low** - The lowest price during the day.
- **Close** - The closing price on the trading day.
- **Adj Close** - The closing price after adjustments on the trading day.
- **Volume** - Number of shares traded.



Importing Libraries and Choosing a Stock from Top 50 American Companies

```
In [1]: #importing all necessary libraries libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.ar_model import AutoReg
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.stattools import adfuller
from pmdarima.arima import auto_arima
from sklearn.metrics import mean_squared_error
from tqdm import tqdm
#import warnings
#warnings.filterwarnings('ignore')

#!pip install yfinance
import yfinance
```

```
In [2]: top50 = pd.read_csv('Top50_American_Companies.csv')
print("----- Choose a Company with its ticker symbol from below list :")
print(top50.head(50))
```

----- Choose a Company with its ticker symbol from below list for Time Series Forecasting -----

Rank (\$B)		Company Name	Ticker Symbol	Market Cap
0	1	APPLE INC.	AAPL	26
10.94				
1	2	MICROSOFT CORPORATION	MSFT	21
27.14				
2	3	ALPHABET INC.	GOOG	13
56.39				
3	4	AMAZON.COM, INC.	AMZN	10
97.20				
4	5	BERKSHIRE HATHAWAY INC.,	BRK.A	7
22.62				
5	6	NVIDIA CORPORATION	NVDA	6
69.84				
6	7	Meta Platforms, Inc.	META	5
46.14				
7	8	TESLA, INC.	TSLA	5
23.19				
8	9	VISA INC.	V	4
81.33				
9	10	EXXON MOBIL CORPORATION	XOM	4
70.92				
10	11	UNITEDHEALTH GROUP INCORPORATED	UNH	4
51.33				
11	12	JOHNSON & JOHNSON	JNJ	4
22.73				
12	13	JPMORGAN CHASE & CO.	JPM	4
11.99				
13	14	WALMART INC.	WMT	4
09.01				
14	15	THE PROCTER & GAMBLE COMPANY	PG	3
68.19				
15	16	ELI LILLY AND COMPANY	LLY	3
66.88				
16	17	MASTERCARD INCORPORATED.	MA	3
57.69				
17	18	CHEVRON CORPORATION	CVX	3
20.42				
18	19	THE HOME DEPOT, INC.	HD	3
03.84				
19	20	MERCK & CO., INC.	MRK	2
92.77				
20	21	ABBVIE INC.	ABBV	2
86.51				
21	22	THE COCA-COLA COMPANY	KO	2
77.10				
22	23	Broadcom Inc.	AVGO	2
63.87				
23	24	ORACLE CORPORATION	ORCL	2
56.89				
24	25	PEPSICO, INC.	PEP	2
55.37				
25	26	BANK OF AMERICA CORPORATION	BAC	2
38.94				

26	27	PFIZER INC.	PFE	2
26.96				
27	28	COSTCO WHOLESALE CORPORATION	COST	2
24.56				
28	29	THERMO FISHER SCIENTIFIC INC.	TMO	2
20.50				
29	30	MCDONALD'S CORPORATION	MCD	2
13.21				
30	31	SALESFORCE, INC.	CRM	1
99.03				
31	32	ABBOTT LABORATORIES	ABT	1
94.06				
32	33	NIKE, INC.	NKE	1
92.94				
33	34	CISCO SYSTEMS, INC.	CSCO	1
92.63				
34	35	DANAHER CORPORATION	DHR	1
84.33				
35	36	THE WALT DISNEY COMPANY	DIS	1
81.90				
36	37	LINDE PUBLIC LIMITED COMPANY	LIN	1
79.72				
37	38	ACCENTURE PUBLIC LIMITED COMPANY	ACN	1
74.51				
38	39	ADOBE INC.	ADBE	1
73.24				
39	40	UNITED PARCEL SERVICE, INC.	UPS	1
67.23				
40	41	TEXAS INSTRUMENTS INCORPORATED	TXN	1
60.62				
41	42	NEXTERA ENERGY, INC.	NEE	1
59.83				
42	43	COMCAST CORPORATION	CMCSA	1
59.11				
43	44	VERIZON COMMUNICATIONS INC.	VZ	1
56.74				
44	45	WELLS FARGO & COMPANY	WFC	1
55.19				
45	46	MORGAN STANLEY	MS	1
52.17				
46	47	Philip Morris International Inc.	PM	1
51.75				
47	48	BRISTOL-MYERS SQUIBB COMPANY	BMJ	1
48.26				
48	49	NETFLIX, INC.	NFLX	1
46.06				
49	50	ADVANCED MICRO DEVICES, INC.	AMD	1
42.32				

```
In [3]: def get_stock_ticker():
        ticker = input("Enter the Selected Company Ticker Symbol: ")
        ticker = str(ticker)
        return ticker

ticker = get_stock_ticker()
```

Enter the Selected Company Ticker Symbol: JPM

```
In [4]: #Retrieving Stock Data for last 10 years
data = yfinance.download(tickers=ticker, start='2018-04-01', end='2023-04-28')
data.head()
```

[*****100%*****] 1 of 1 completed

Out[4]:

	Open	High	Low	Close	Adj Close	Volume
Date						
2018-04-02	109.959999	110.730003	106.080002	107.849998	92.371223	18822500
2018-04-03	108.360001	109.500000	107.260002	109.330002	93.638824	14050700
2018-04-04	107.099998	111.209999	107.019997	110.989998	95.060562	15302600
2018-04-05	111.629997	112.830002	111.389999	111.879997	96.308762	16627000
2018-04-06	110.550003	111.550003	107.820000	109.089996	93.907066	18906000

```
In [5]: #Shape of the Data
data.shape
```

Out[5]: (1279, 6)

```
In [6]: #Getting information about null values and data types
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1279 entries, 2018-04-02 to 2023-04-28
Data columns (total 6 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   Open          1279 non-null   float64
 1   High          1279 non-null   float64
 2   Low           1279 non-null   float64
 3   Close         1279 non-null   float64
 4   Adj Close     1279 non-null   float64
 5   Volume        1279 non-null   int64   
dtypes: float64(5), int64(1)
memory usage: 69.9 KB
```

```
In [7]: data.index = pd.DatetimeIndex(data.index)
```

```
In [8]: #Add returns to the dataframe
Returns = data.pct_change().dropna()
Returns
```

Out[8]:

	Open	High	Low	Close	Adj Close	Volume
Date						
2018-04-03	-0.014551	-0.011108	0.011124	0.013723	0.013723	-0.253516
2018-04-04	-0.011628	0.015616	-0.002238	0.015183	0.015183	0.089099
2018-04-05	0.042297	0.014567	0.040834	0.008019	0.013131	0.086547
2018-04-06	-0.009675	-0.011344	-0.032050	-0.024937	-0.024937	0.137066
2018-04-09	-0.004885	0.012819	0.019199	0.012008	0.012008	-0.155067
...
2023-04-24	0.005152	-0.000071	0.009079	0.001352	0.001352	-0.360988
2023-04-25	-0.006123	-0.008079	-0.017209	-0.021744	-0.021744	0.404599
2023-04-26	-0.014183	-0.015362	-0.026375	-0.017724	-0.017724	0.356261
2023-04-27	-0.011772	0.000726	0.012388	0.013459	0.013459	-0.372847
2023-04-28	0.004118	0.002973	0.000074	0.008683	0.008683	0.092986

1278 rows × 6 columns

Descriptive Statistics

```
In [9]: data.describe().T
```

Out[9]:

	count	mean	std	min	25%	50%	
Open	1279.0	1.242790e+02	2.136728e+01	8.156000e+01	1.081700e+02	1.178800e+02	1.389100e+02
High	1279.0	1.255491e+02	2.140365e+01	8.375000e+01	1.092400e+02	1.189000e+02	1.401300e+02
Low	1279.0	1.229943e+02	2.131766e+01	7.691000e+01	1.070250e+02	1.163200e+02	1.377900e+02
Close	1279.0	1.242686e+02	2.137137e+01	7.903000e+01	1.080300e+02	1.178500e+02	1.390800e+02
Adj Close	1279.0	1.154550e+02	2.289772e+01	7.148599e+01	9.503904e+01	1.119901e+02	1.348300e+02
Volume	1279.0	1.440600e+07	6.845342e+06	3.220500e+06	1.001935e+07	1.268180e+07	1.635700e+07


```
In [10]: print("Minimum Adjusted Close recorded in last 10 years: ", round(data[
print("\nMaximum Adjusted Close recorded in last 10 years: ", round(data[
per_inc = (data['Adj Close'].iloc[data.shape[0]-1]-data['Adj Close'].iloc[0])/data['Adj Close'].iloc[0]
print("\nPercentage Increase for Adjusted Close recorded in last 10 years: ", round(per_inc*100, 2))
```

Minimum Adjusted Close recorded in last 10 years: 71.486

Maximum Adjusted Close recorded in last 10 years: 164.016

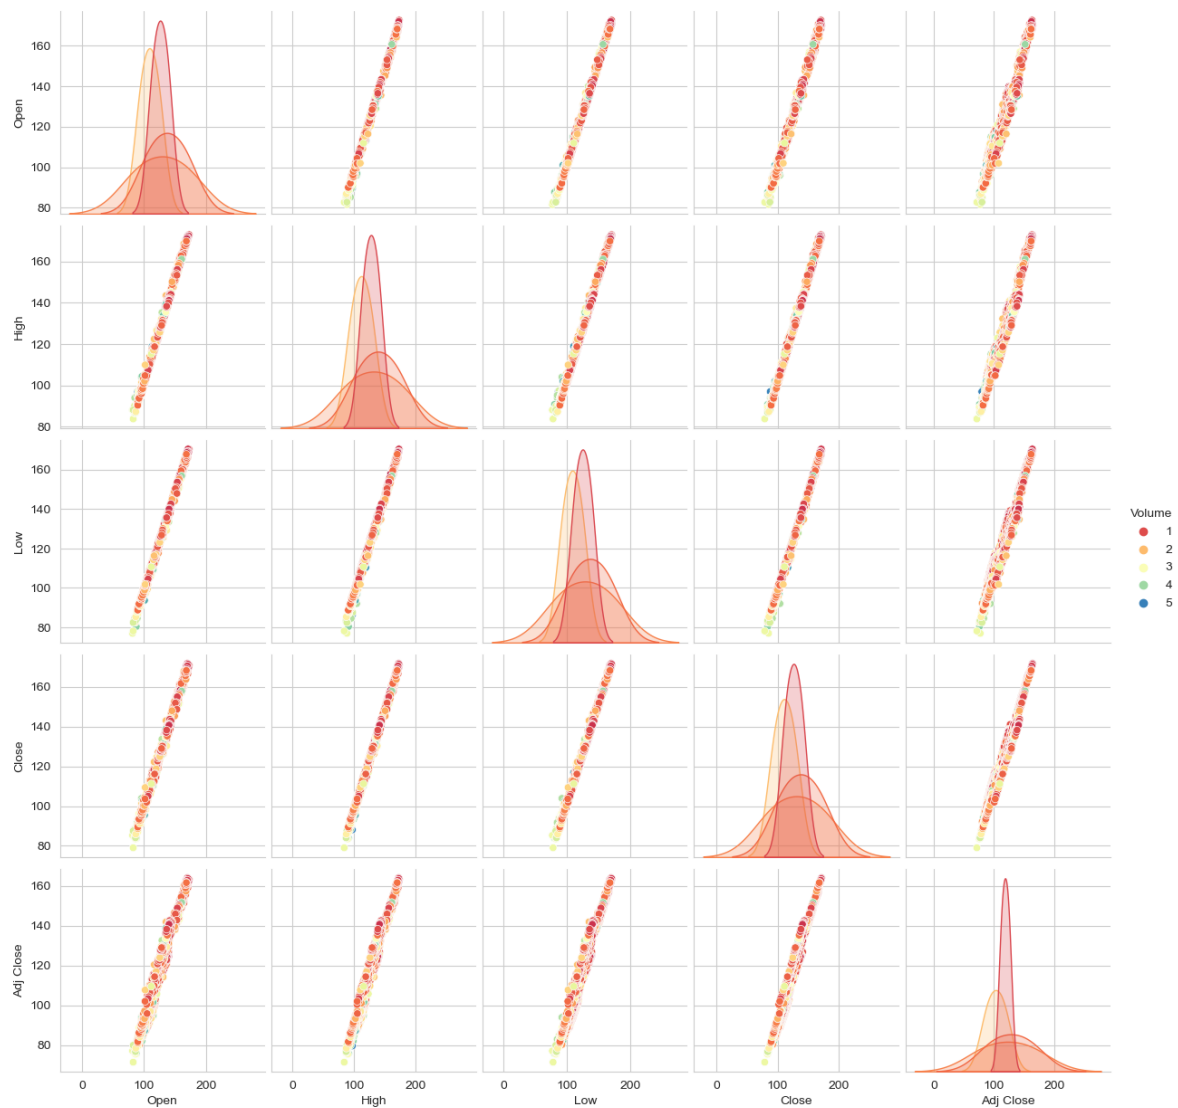
Percentage Increase for Adjusted Close recorded in last 10 years (4/2013 vs 3/2023): 49.66%

```
In [11]: #Visualizing box plot for Adjusted Close column
sns.boxplot(y=data['Close'], color='grey')
plt.show()
```

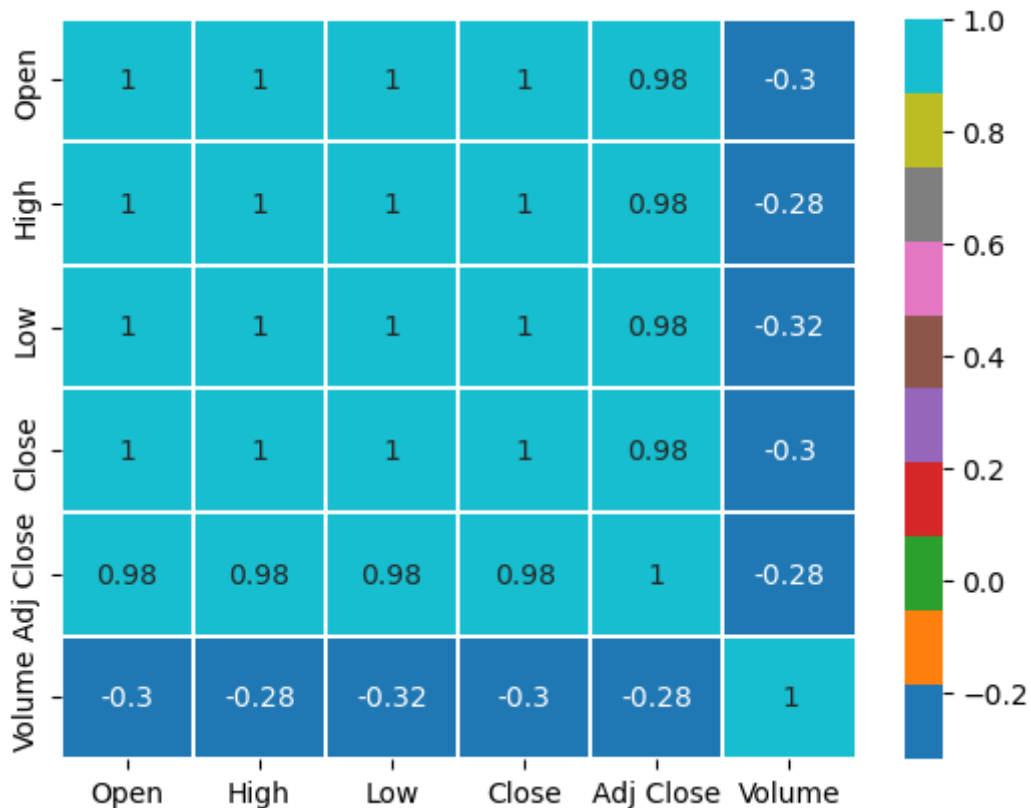



```
In [48]: sns.pairplot(data, hue='Volume', palette='Spectral')
```

```
Out[48]: <seaborn.axisgrid.PairGrid at 0x7f91ccbd2d60>
```



```
In [12]: #finding correlation for the data
c=data.corr()
sns.heatmap(c, cmap='tab10', linewidths='0.05', annot=True)
plt.show()
```




Except Volume all other columns are in perfect positive correlation with each other

Moving Averages for this Data

The moving average is also known as the rolling mean and is calculated by averaging data of the time series within k periods of time.

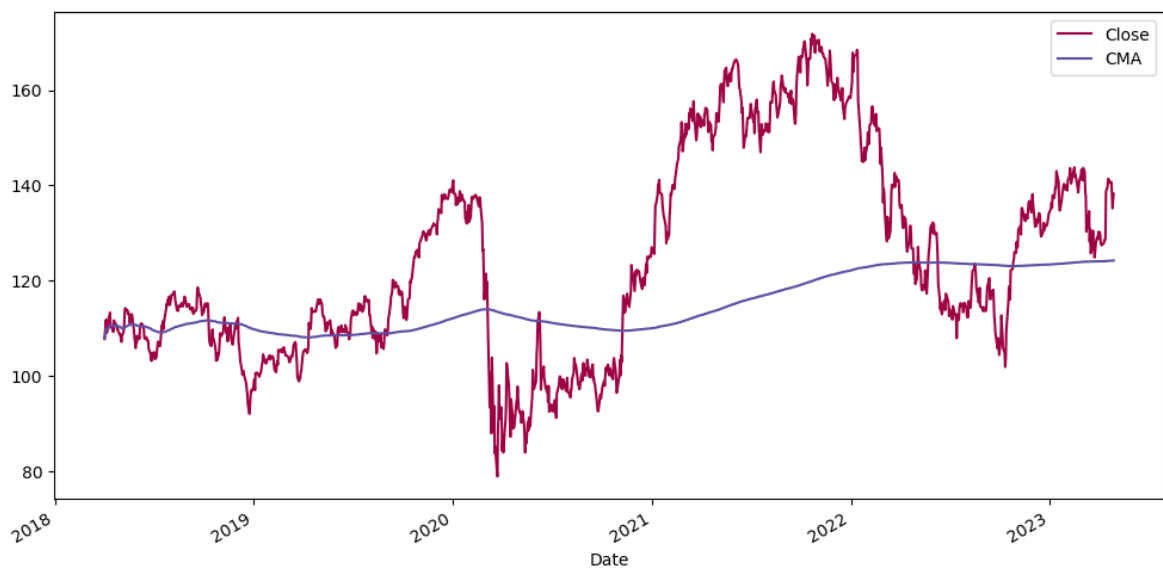
1. Simple Moving Average (SMA)

```
In [13]: def Simple_Moving_Avg(df, feature, num_of_observations):  
  
    #Using .to_frame to convert the column to Pandas.Series  
    new_df = df[feature].to_frame()  
  
    #Using rolling(window).mean() to calculate SMA  
    #using window size i.e, num of observations  
    new_df['SMA'] = new_df[feature].rolling(num_of_observations).mean()  
  
    #plotting SMA  
    new_df.plot(figsize=(12,6), colormap='coolwarm')  
    plt.show()  
Simple_Moving_Avg(data, 'Close', 60)
```




2. Cumulative Moving Average (CMA)

```
In [14]: def Cum_Moving_Avg(df, feature):  
  
    #Using .to_frame to convert the column to Pandas.Series  
    new_df = df[feature].to_frame()  
  
    #Using expanding().mean() to calculate CMA  
    new_df['CMA'] = new_df[feature].expanding().mean()  
  
    #plotting SMA  
    new_df.plot(figsize=(12,6), colormap='Spectral')  
    plt.show()  
Cum_Moving_Avg(data, 'Close')
```



3. Exponential Moving Average (EMA)

```
In [15]: def Expo_Moving_Avg(df, feature, num_of_observations):  
  
    #Using .to_frame to convert the column to Pandas.Series  
    new_df = df[feature].to_frame()  
  
    #Using ewm(span).mean() to calculate EMA  
    #using span i.e, num of observations  
    new_df['EMA'] = new_df[feature].ewm(span=num_of_observations).mean()  
  
    #plotting SMA  
    new_df.plot(figsize=(12,6))  
    plt.show()  
Expo_Moving_Avg(data, 'Close', 60)
```



```
In [16]: #Trends for priced variables
sns.set_style('whitegrid')
price_features = (data.drop(['Volume'], axis=1)).columns
data[price_features].plot(figsize=(12,6), colormap='Spectral')
plt.show()
```



Time Series Decomposition

```
In [17]: #Applying time series decomposition
def decompose_results(data, feature, model, period=180):


    df = data
    df.reset_index(inplace=True)
    df.set_index('Date', inplace=True)
    df = df[feature]

    if model == 'additive':
        decompose_add = seasonal_decompose(df, model=model, period=period)
        plt.figure(figsize=(16,4))
        decompose_add.plot()
        plt.xlabel('Additive', fontsize=18)
        plt.show()

    elif model == 'multiplicative':
        decompose_mul = seasonal_decompose(df, model=model, period=period)
        plt.figure(figsize=(16,4))
        decompose_mul.plot()
        plt.xlabel('Multiplicative', fontsize=18)
        plt.show()
```

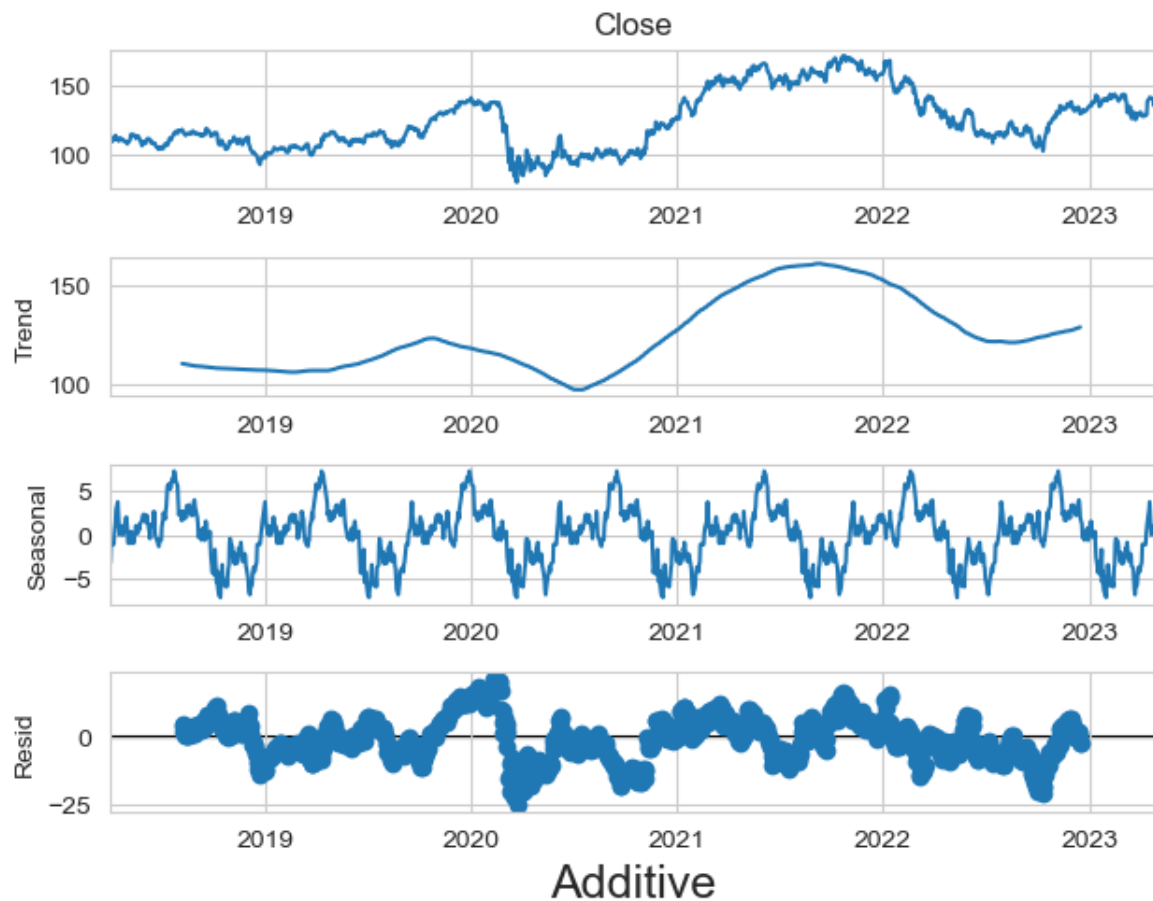
```
In [18]: #Seasonal Decomposition with Open column  
decompose_results(data, 'Close', 'multiplicative')
```

<Figure size 1600x400 with 0 Axes>



```
In [19]: #Seasonal Decomposition with Close column  
decompose_results(data, 'Close', 'additive')
```


<Figure size 1600x400 with 0 Axes>



Auto Correlation using lag_plot() and plot_acf() functions

```
In [20]: #lag_plot()  
pd.plotting.lag_plot>Returns, lag=3)
```

```
Out[20]: <AxesSubplot:xlabel='y(t)', ylabel='y(t + 3) '>
```



```




In [21]: #q=2
#Plotting Auto Correlation using plt_acf()
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16,4))

ax1.plot(data['Close'])
ax1.set_title('Original')
plot_acf(data['Close'], ax=ax2)
plt.show()

fig, (ax3, ax4) = plt.subplots(1, 2, figsize=(16,4))

ax3.plot>Returns['Close'])
ax3.set_title('Original')
plot_acf>Returns['Close'], ax=ax4)
ax4.set_ylim(0,1)
plt.show()

```



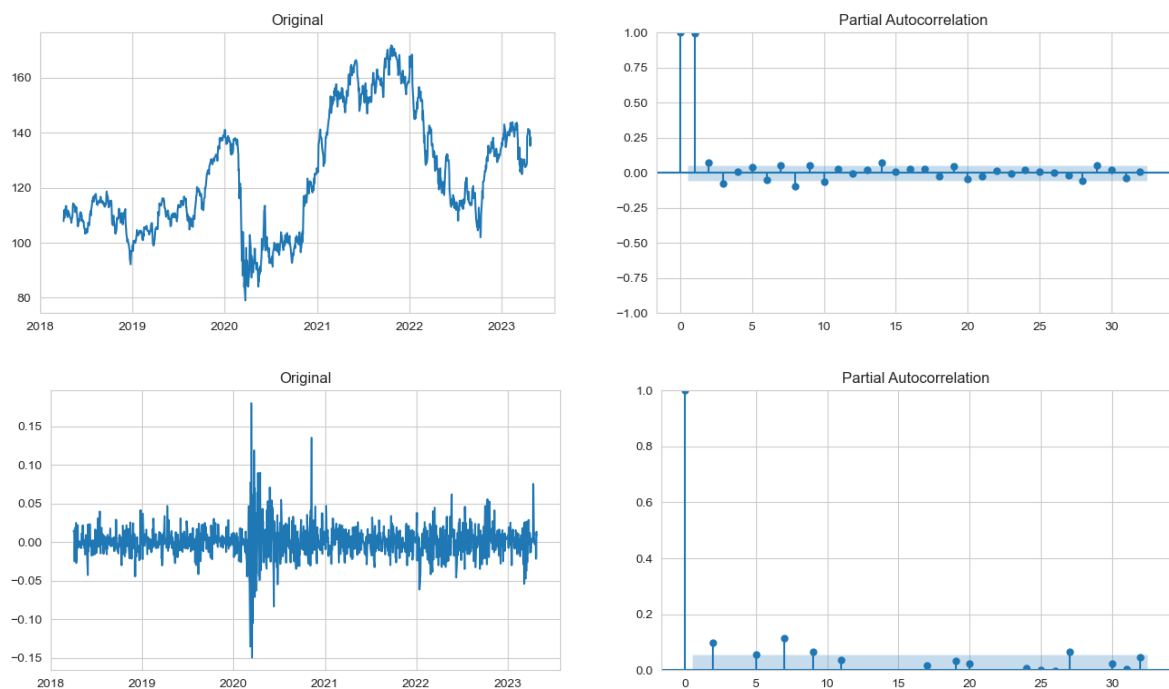
Note: The partial autocorrelation function, like the ACF, indicates only the association between two data that the shorter lags between those observations do not explain.

```
In [22]: #p=7
#Plotting Auto Correlation using plt_pacf()
fig, (ax5, ax6) = plt.subplots(1, 2, figsize=(16,4))

ax5.plot(data['Close'])
ax5.set_title('Original')
plot_pacf(data['Close'], ax=ax6, method='ywmm')
plt.show()

fig, (ax7, ax8) = plt.subplots(1, 2, figsize=(16,4))

ax7.plot>Returns['Close'])
ax7.set_title('Original')
plot_pacf>Returns['Close'], ax=ax8, method='ywmm')
ax8.set_ylim(0,1)
plt.show()
```



We can use pmdarima to get the number of differencing and ndiffs function is used to estimate ARIMA differencing term 'd'

```
In [23]: from pmdarima.arma.utils import ndiffs
ndiffs(data.Close, test='adf')
```

Out[23]: 1

Augumented Dickey - Fuller Test

The Augmented Dickey-Fuller Test is used to determine if time-series data is stationary or not.


```
In [24]: adftest = adfuller(data['Close'], autolag='AIC')
adfoutput = pd.Series(adftest[0:4], index=['Test Statistic', 'p-value', '#Lags Used', 'Number of Observations Used'])
for key, value in adftest[4].items():
    adfoutput['Critical Value (%s)' % key] = value
print(adfoutput)
```

```
Test Statistic          -1.796593
p-value                  0.382154
#Lags Used              13.000000
Number of Observations Used 1265.000000
Critical Value (1%)     -3.435530
Critical Value (5%)     -2.863827
Critical Value (10%)    -2.567988
dtype: float64
```


So, the data is not stationary as we can observe that p-value is higher than significance value.

```
In [25]: from scipy.stats import boxcox
data_boxcox = pd.Series(boxcox(data['Close'], lmbda=0), index = data.index)

plt.figure(figsize=(12,4))
plt.plot(data_boxcox, label='After Box Cox tranformation')
plt.legend(loc='best')
plt.title('After Box Cox transform')
plt.show()
```



```
In [26]: data_boxcox_diff = pd.Series(data_boxcox - data_boxcox.shift(), data.index)
plt.figure(figsize=(12,4))
plt.plot(data_boxcox_diff, label='After Box Cox tranformation and differencing')
plt.legend(loc='best')
plt.title('After Box Cox transform and differencing')
plt.show()
```



```
In [27]: data_boxcox_diff.dropna(inplace=True)
```

```
In [28]: adf_test1 = adfuller(data_boxcox_diff)

print('ADF Statistic: %f' % adf_test1[0])
print('Critical Values @ 0.05: %.2f' % adf_test1[4]['5%'])
print('p-value: %f' % adf_test1[1])
```

```
ADF Statistic: -9.632163
Critical Values @ 0.05: -2.86
p-value: 0.000000
```

Train_Test_Split

```
In [29]: train = data[data.index < pd.to_datetime("2022-01-01", format='%Y-%m-%d')
test = data[data.index > pd.to_datetime("2022-01-01", format='%Y-%m-%d')

plt.plot(train.index, train['Close'], color = "black")
plt.plot(test.index, test['Close'], color = "blue")
plt.ylabel(ticker)
plt.xlabel('Date')
plt.xticks(rotation=45)
plt.title("Train/Test Split")
plt.show()
```




Holt Winter's multiplicative method with trend and seasonality

```
In [30]: from statsmodels.tsa.holtwinters import ExponentialSmoothing
hwm = test.copy()
model = ExponentialSmoothing(np.asarray(train['Close']), seasonal_period=12,
                             model='ML', damped_trend=True)
model_fit = model.fit(optimized=True)
print(model_fit.params)
hwm['forecast'] = model_fit.forecast(len(test))
```

```
{'smoothing_level': 0.8866541673260788, 'smoothing_trend': 0.0, 'smoothing_seasonal': 0.0, 'damping_trend': nan, 'initial_level': 108.55499521212437, 'initial_trend': 0.05271142113027826, 'initial_seasons': array([0.99490971, 0.99313605, 0.99227405, 0.99391657, 0.9928215, 0.99534114, 0.99644682, 0.99768944, 1.00101302, 0.99987383, 1.00020544, 1.000072]), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```

```
In [31]: #pred = test.copy()
#y_hat_holt['holt_forecast'] = model_fit.forecast(len(test))
plt.figure(figsize=(12,4))
plt.plot(train['Close'], label='Train')
plt.plot(test['Close'], label='Test')
plt.plot(hwm['forecast'], label='Holt\'s Exponential Smoothing Forecast')
plt.legend(loc='best')
plt.title('Holt\'s Exponential Smoothing Method')
plt.show()
```



```
In [32]: rmse = np.sqrt(mean_squared_error(test['Close'], hwm['forecast'][test.index]))
mape = np.round(np.mean(np.abs(test['Close']-hwm['forecast'][test.index])/test['Close']), 2)

Results_hwm = pd.DataFrame({'Method': ['Holt Winter\'s multiplicative method'], 'RMSE': rmse, 'MAPE': mape})
Results_hwm
```

Out[32]:

	Method	RMSE	MAPE
0	Holt Winter's multiplicative method with trend...	39.349	29.598

ARIMA model

```
In [33]: #resolving indices to ensure error free modelling
d=data.reset_index()
data.index=d['Date']

tr=train.reset_index()
train.index=tr['Date']

te=test.reset_index()
test.index=te['Date']
```



```
In [34]: ar_model = ARIMA(test['Close'], order=(7,1,2))
ar_model_fit = ar_model.fit()
print(ar_model_fit.summary())
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
```

```
    self._init_dates(dates, freq)
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
```

```
    self._init_dates(dates, freq)
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
```

```
    self._init_dates(dates, freq)
```

SARIMAX Results

```

=====
=====
Dep. Variable:          Close    No. Observations:
332
Model:                ARIMA(7, 1, 2)    Log Likelihood
-754.452
Date:                Tue, 09 May 2023    AIC
1528.904
Time:                22:31:49    BIC
1566.925
Sample:                0    HQIC
1544.069

- 332
Covariance Type:      opg
=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.1195	0.258	0.463	0.643	-0.386
0.625					
ar.L2	0.7881	0.197	3.991	0.000	0.401
1.175					
ar.L3	-0.0470	0.072	-0.654	0.513	-0.188
0.094					
ar.L4	0.0728	0.067	1.080	0.280	-0.059
0.205					
ar.L5	-0.0300	0.076	-0.394	0.694	-0.180
0.119					
ar.L6	-0.1592	0.059	-2.702	0.007	-0.275
-0.044					
ar.L7	0.0552	0.071	0.775	0.438	-0.085
0.195					
ma.L1	-0.0561	0.252	-0.223	0.824	-0.550
0.438					
ma.L2	-0.7788	0.192	-4.047	0.000	-1.156
-0.402					
sigma2	5.5849	0.352	15.884	0.000	4.896
6.274					

```

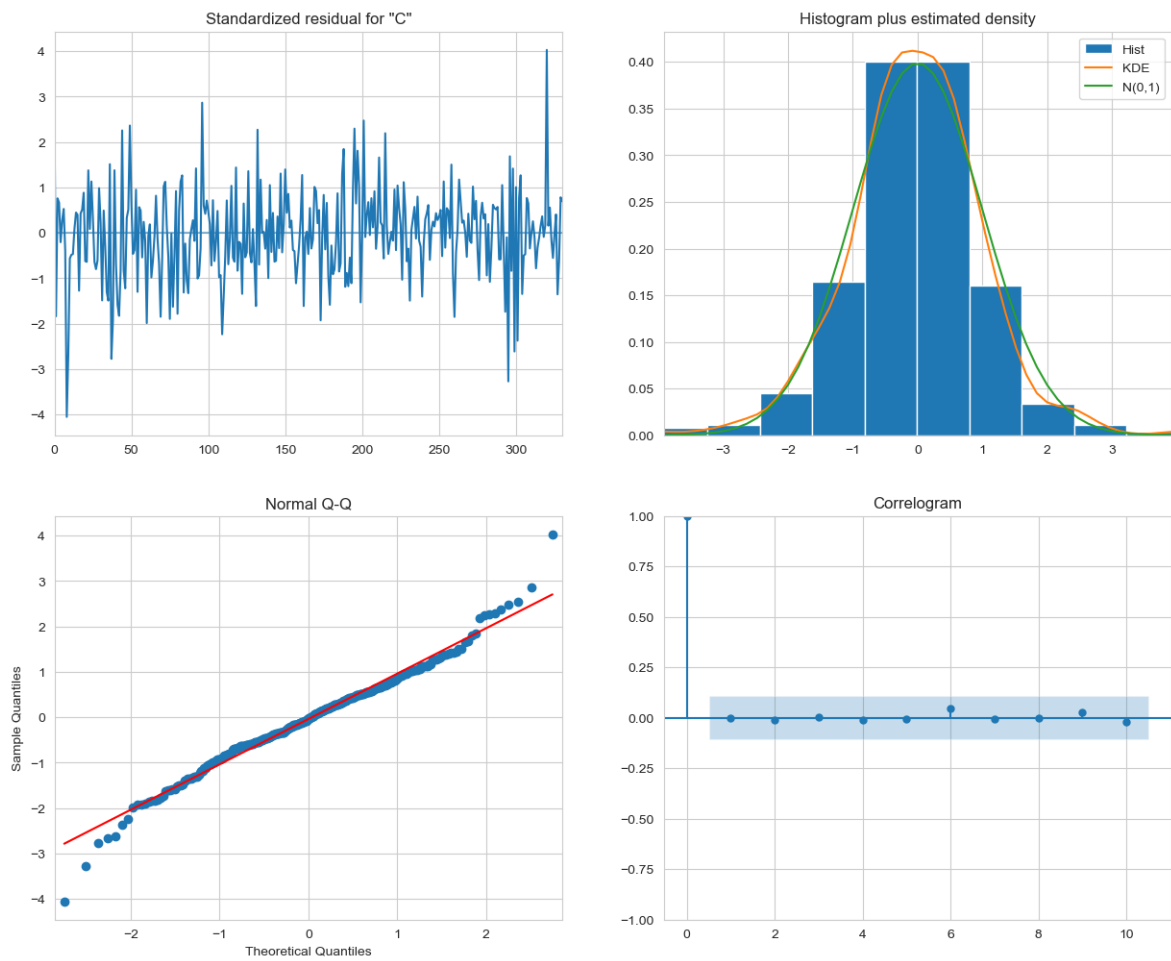
=====
=====
Ljung-Box (L1) (Q):          0.00    Jarque-Bera (JB):
33.06
Prob(Q):                    1.00    Prob(JB):
0.00
Heteroskedasticity (H):      0.64    Skew:
-0.08
Prob(H) (two-sided):        0.02    Kurtosis:
4.54
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [35]: ar_model_fit.plot_diagnostics(figsize=(15, 12))  
plt.show()
```



```
In [36]: preds = ar_model_fit.predict()
preds[3:].plot()
test['Close'].plot()
plt.legend(loc='best')
```

Out[36]: <matplotlib.legend.Legend at 0x7f91a87e4d90>



```
In [37]: rmse = np.sqrt(mean_squared_error(test['Close'], preds[test.index.min():]))
mape = np.round(np.mean(np.abs(test['Close']-preds[test.index.min():])/test['Close']), 2)

Results_ar = pd.DataFrame({'Method': ['Autoregressive Integrated Moving Average (ARIMA)'],
                           'RMSE': rmse,
                           'MAPE': mape})
Results_ar
```

Out[37]:

	Method	RMSE	MAPE
0	Autoregressive Integrated Moving Average (ARIMA)	9.183	1.695

Optimization of the Model

Adjusting parameters within your modeling function (#p - lagdays to use, #d - degree of differencing, #q - degree of seasonality)

AutoRegression Method (only #p - lagdays to use) using ARIMA

```
In [38]: #After some trials adjusting p,d,q, I have decided to use AutoRegression  
areg_model = ARIMA(test['Close'], order=(7,0,0))  
areg_model_fit = areg_model.fit()  
print(areg_model_fit.summary())
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_m  
odel.py:471: ValueWarning: A date index has been provided, but it has  
no associated frequency information and so will be ignored when e.g.  
forecasting.
```

```
    self._init_dates(dates, freq)
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_m  
odel.py:471: ValueWarning: A date index has been provided, but it has  
no associated frequency information and so will be ignored when e.g.  
forecasting.
```

```
    self._init_dates(dates, freq)
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_m  
odel.py:471: ValueWarning: A date index has been provided, but it has  
no associated frequency information and so will be ignored when e.g.  
forecasting.
```

```
    self._init_dates(dates, freq)
```

SARIMAX Results

```

=====
=====
Dep. Variable:          Close    No. Observations:
332
Model:                ARIMA(7, 0, 0)    Log Likelihood
-760.477
Date:                Tue, 09 May 2023    AIC
1538.954
Time:                22:31:50    BIC
1573.200
Sample:                0    HQIC
1552.611
                        - 332
Covariance Type:      opg
=====
=====
                        coef    std err          z      P>|z|      [0.025
0.975]
-----
const          135.0078      6.365      21.212      0.000      122.533
147.483
ar.L1           1.0496      0.051      20.739      0.000       0.950
1.149
ar.L2          -0.0334      0.077      -0.437      0.662      -0.184
0.117
ar.L3          -0.0197      0.075      -0.261      0.794      -0.167
0.128
ar.L4           0.0840      0.077       1.087      0.277      -0.067
0.235
ar.L5          -0.1116      0.082     -1.363      0.173      -0.272
0.049
ar.L6          -0.0187      0.081     -0.231      0.817      -0.177
0.140
ar.L7           0.0348      0.058       0.601      0.548      -0.079
0.148
sigma2          5.6522      0.360     15.700      0.000       4.947
6.358
=====
=====
Ljung-Box (L1) (Q):          0.00    Jarque-Bera (JB):
28.20
Prob(Q):          0.96    Prob(JB):
0.00
Heteroskedasticity (H):      0.61    Skew:
0.01
Prob(H) (two-sided):        0.01    Kurtosis:
4.43
=====
=====

```

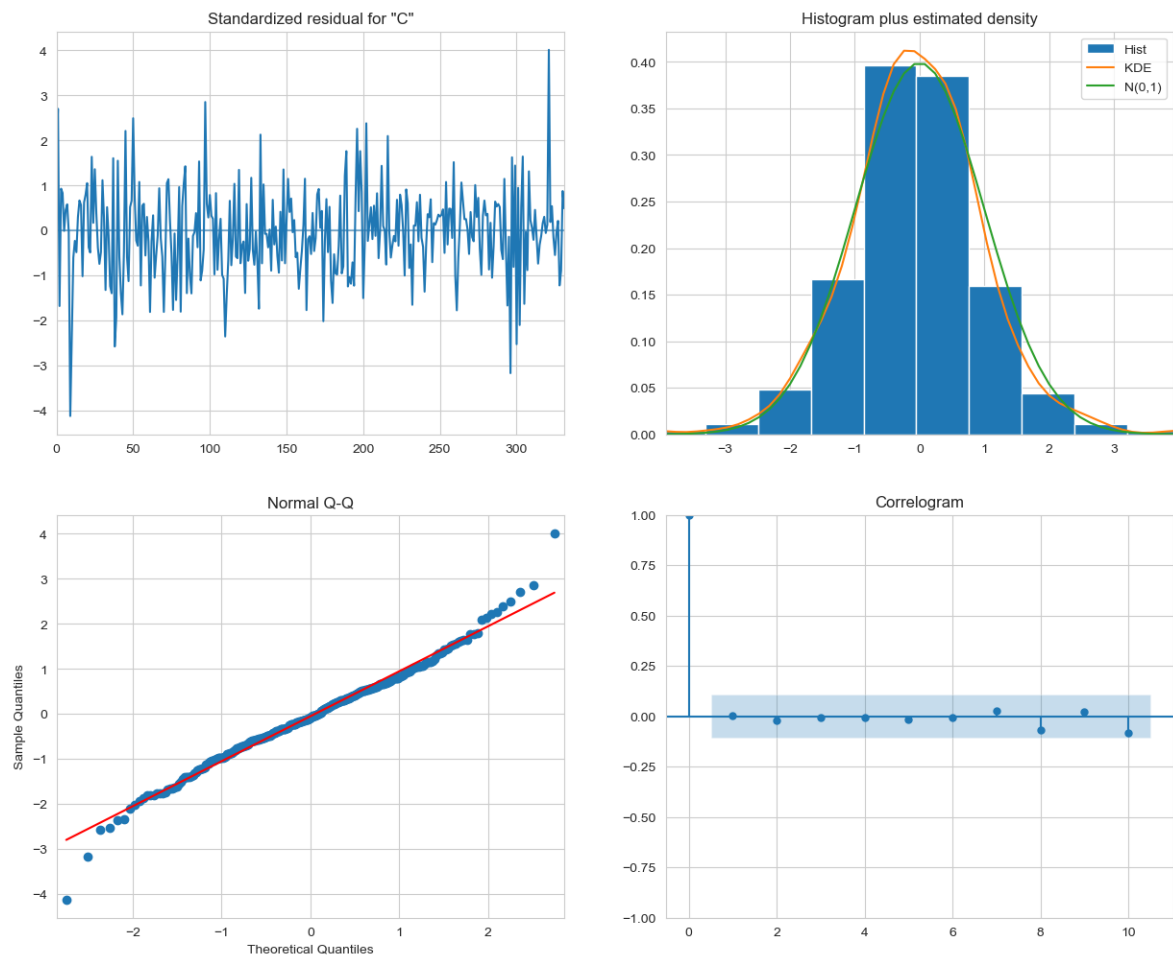
Warnings:

```

[1] Covariance matrix calculated using the outer product of gradients
(complex-step).


```

```
In [39]: #plotting Diagnostics  
areg_model_fit.plot_diagnostics(figsize=(15, 12))  
plt.show()
```



```
In [40]: preds_areg = areg_model_fit.predict()
preds_areg[3:].plot()
test['Close'].plot()
plt.legend(loc='best')
```

Out[40]: <matplotlib.legend.Legend at 0x7f91c97d1e20>



```
In [41]: rmse = np.sqrt(mean_squared_error(test['Close'], preds_areg[test.index.min():]))
mape = np.round(np.mean(np.abs(test['Close'] - preds_areg[test.index.min():])), 3)

Results_areg = pd.DataFrame({'Method': ['Autoregression (AR) method'], 'RMSE': rmse, 'MAPE': mape})
Results_areg
```

Out[41]:

	Method	RMSE	MAPE
0	Autoregression (AR) method	2.783	1.455


```
In [42]: #Presenting the results
results = pd.concat([Results_hwm, Results_ar, Results_areg])
results = results[['Method', 'RMSE', 'MAPE']]
results
```

Out[42]:

	Method	RMSE	MAPE
0	Holt Winter's multiplicative method with trend...	39.349	29.598
0	Autoregressive Integrated Moving Average (ARIM...	9.183	1.695
0	Autoregression (AR) method	2.783	1.455

Summary:

Note: This Time Series Forecasting Project is generalized one which can accept any Stock data with its ticker which is present from 2018 to present.

1. Considered Stock data among the top 50 American Companies which is JPMORGAN CHASE & CO.(Ticker='JPM').
2. Stock data is non-null data containing Date as its index.
3. The Inter-Quartile Range for Close column is 110-140 (approx).
4. There is considerable variation for Adjusted Close variable per Volume as compared to others.
5. Exponential Moving Average shows better fit than Cumulative and Simple Moving Averages.
6. Adjusted Close which has variation with other price columns earlier moved closer in recent years.
7. Additive decomposition showing high noise(resid) variability.
8. Seasonality is observed with this stock data.
9. Degree of Seasonality(q) is considered as 2 by observing AutoCorrelation Plot.
10. Number of lags(p) can be considered 2 or 7 as there are spikes above the significance region for Partial AutoCorrelation plot.
11. After performing Augmented Dickey-Fuller test; p-value>0.05, So, the data is not stationary. Also boxcox transform to make it Stationary and its differencing is same as the plot for Returns Dataframe with Close variable.
12. Holt Winter's Multiplicative method with trend and seasonality has higher RMSE than ARIMA model.
13. Histogram for residuals in plot_diagnostics has higher maximum density than Optimized model.
14. Mean Absolute Percentage Error reduced well between Holts Winter's method and ARIMA model.
15. Optimization of ARIMA Model is Autoregression model with only lags used and is the best model with low Root-Mean-Squared-Error.

- Jagadesh Varma Nadimpalli